

Stat 628: Data Science Practicum

Spring 2018, UW-Madison

Module 2 Goals and Data Background

Motivation and Goal:

Consider the following reviews from Yelpers reviewing the restaurant 1847 At the Stamm House in Madison.

“We've been waiting for 1847 Stamm House to open. Glad to miss the opening kinks being worked out some reviews mentioned. Went last night, and it was fabulous!! The fish fry was a new take on a Wisconsin fav, a whole piece of whitefish with delicious cole slaw and picked sides. All amazing! Can't wait to return to try the rest of the menu. Maybe on a night that's quiet, if there is one. Glad to see the place was hopping!”

“I was visiting from out of state and had this place recommended for a Friday night fish fry. What a mistake. I was told over the phone that the Cod fish fry was \$15. I person we were told the Cod cost \$16 and the Bluegill was \$18. When the bill came I was informed that the Cod was \$18. For one lousy piece of Cod. Outrageously overpriced. I also ordered a Moscow Mule. That was served in a glass, instead of the customary copper mug.”

“I too am hoping that this unique and beautiful restaurant works out the kinks and begins to perform up to the high expectations. My wife and I experienced the same problems during week #1 that the other reviewers have described (e.g. the relatively expensive burger and not-so-good white fish). We loved the atmosphere, albeit a bit too noisy, but the food and the service didn't measure up to the competition that day. I suggest the chef and servers visit Brasserie V to experience another way of doing things! Their best burger is actually worth \$12! We are looking forward to our next meal at the Stamm House and are sure we will experience a marked improvement in all aspects of the dining experience!”

It's easy for humans to understand the *sentiment* behind the reviews, specifically whether a review is positive or negative, even without having any context about the business. In fact, it's pretty easy for a human to guess the Yelp ratings based on the text alone. In contrast, for machines, understanding the sentiment of a text is difficult because it requires understanding not only the underlying language structure, but also the meaning it conveys. In fact, this is an active area of research in a field known as “sentiment analysis.”

The two goals of this project are: (1) find out **what** makes a review positive or negative based on the review and a small set of attributes and (2) propose a prediction model to **predict the ratings of reviews** based on the **text and said attributes** and submit the predictions to Kaggle (see grading description document).

Data Background:

Yelp is an Internet company founded in 2004 to “help people find great local businesses” by providing a platform for users to write reviews of businesses. As of Q4 2016, there are more than 121 million reviews and 89 million unique (average) monthly visitors have visited its website. Yelp utilizes automated software (https://www.yelp-support.com/article/What-is-Yelp-s-recommendation-software?l=en_US) to filter out bad reviews.

Recently, Yelp has released some of its 121 million reviews to the public as a part of the “Yelp Dataset Challenge.” This data contains 4.1 million reviews from 144K businesses from the following cities: Edinburgh (U.K.), Karlsruhe (Germany), Montreal (Canada), Waterloo (Canada), Pittsburgh (U.S.), Charlotte (U.S.), Urbana-Champaign (U.S.), Phoenix (U.S.), Las Vegas (U.S.), Madison (U.S.), Cleveland (U.S.). Yelp “challenges students to use this data in an innovative way and break ground in research” and they provide cash rewards for the winners.

Our class project will focus one area of natural language processing, called **sentiment analysis**, specifically (i) finding out what makes a review positive or negative and (ii) predicting a review’s rating based on its text and a small set of relevant attributes; to make the data more manageable on my hard-drive, I had to take out some specific attributes and focus on a subset of Yelp reviews. Specifically, our subset consists of reviews with the tag “Restaurants” as part of its category values. The businesses that are included in the data must have at least 3 reviews older than 14 days. Also, only reviews that were recommended at the time of the data collection are included (see link here for how Yelp recommends reviews to its users https://www.yelp-support.com/Recommended_Reviews).

In total, you will be analyzing about 1.5 million reviews with features whose dimension is truly high dimensional (i.e. text data).

Training, Testing, and Validation Data

We have randomly split the data into three pieces whereby 60% of the data is training data, 20% of the data is testing data, and the last 20% is validation data. You’ll have two files, one of the training data and the other that combines the testing + validation data; the identity of the testing/validation within the combined data will not be revealed.

The link to the training data file is here:

https://drive.google.com/open?id=1X5zUFW_elTTBs9ky8d_VfbkEsxs0qgEk

The link to the combined testing and validation data file is here

https://drive.google.com/open?id=1DtBDJ0bRVeo8WG_BXf_ZF7jCYMeXl7Em

You will need to sign in with your UW-Madison e-mail when the Google Log-In shows up. Then, it will redirect you to the UW-Madison login page and you should be able to download the data.

Variable Names and Predictors

We include the following predictors for you to use in your data analysis.

Name	Description
stars	Stars of the review (1=worst, 5=best). The value only takes integers between 1 and 5. This is only available in the training data.
id	ID number for Kaggle. This is only available for the test and validation data.
name	Name of the business
text	Review text
date	Date of review
city	City in which the business is located
longitude	Longitude coordinates of the business's location
latitude	Latitude coordinates of the business's location
categories	An un-parsed string that categorizes the businesses into Yelp's categories. See the full list here: https://www.yelp.com/developers/documentation/v2/category_list

As you can see, the data is not cleaned, especially the text and the categories portion. You are **strongly encouraged** to **come up with new predictors** based on the predictors we have given you, especially the review text. HOWEVER, you must follow the guidelines set below; failure to follow these guidelines will result in an **automatic F for the course and other school-level disciplinary actions**.

Rules and Academic Integrity

Each student assumes the responsibilities of an active participant in UW-Madison's community of scholars in which everyone's academic work and behavior are held to the highest academic integrity standards. Academic misconduct compromises the integrity of the university. Cheating, fabrication, plagiarism, sabotaging other groups' work, unauthorized collaboration, and helping others commit these acts are examples of academic misconduct. Specific examples include, but are not limited to,

1. Copying, plagiarizing, stealing, fabricating any of the deliverables, especially the code or the predictions on Kaggle, from other groups, students outside of the class, or the Internet. In particular, while you may ask other groups for general ideas and questions, you cannot ask for help cleaning the data set, analyzing the dataset, and doing other activities that would be inconsistent with the academic integrity at UW-Madison. If you are unsure, you are always welcome to ask the TA or the professor.
2. Using unauthorized sources, including the original Yelp dataset on Yelp's website or the original ratings (or summaries of ratings of businesses) which can be derived from Yelp's website. You are also not allowed to directly copy, steal, plagiarize, paraphrase, or use any analysis that was already conducted on the Yelp data by others (e.g. data science courses online, someone's blog post or R markdown, Google Cloud's API platform for sentiment analysis, any pre-written software/code that does sentiment analysis automatically, etc.).

However, you are **strongly encouraged** to browse through Yelp, resources on natural language processing (NLP), sentiment analysis, and other researchers' analysis of the Yelp data and gather **background information**. You are strongly encouraged to use the information from your background research **to complement** your own analysis and **provide proper attributions**. In short, your analysis of the data must be **original** and **must be your own work**. Or, in industry-lingo, you should not be stealing others' intellectual property.

If you have any questions about this, please come talk to the TA or the professor.

3. Attempting to gain an unfair advantage by recreating the original Yelp data and using predictors that are not part of the data set. You must only work with the data set you were provided with.

You are strongly encouraged to create your own predictors based on the data set you were given. Again, please come talk to the TA or the professor if you have any questions about this.

4. You may not ask someone to do any part of the analysis on your behalf.

Committing said acts can result in disciplinary action, which includes, but is not limited to failure on the assignment/course, disciplinary probation, or suspension. Substantial or repeated cases of misconduct will be forwarded to the Office of Student Conduct & Community Standards for additional review. For more information, refer to students.wisc.edu/student-conduct/academic-integrity/.